

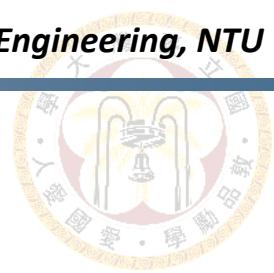
Computer-Aided VLSI System Design

Homework 3: Simple Convolution and Image Processing Engine

Graduate Institute of Electronics Engineering, National Taiwan University

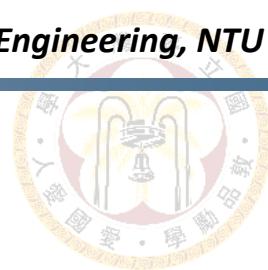


NTU GIEE



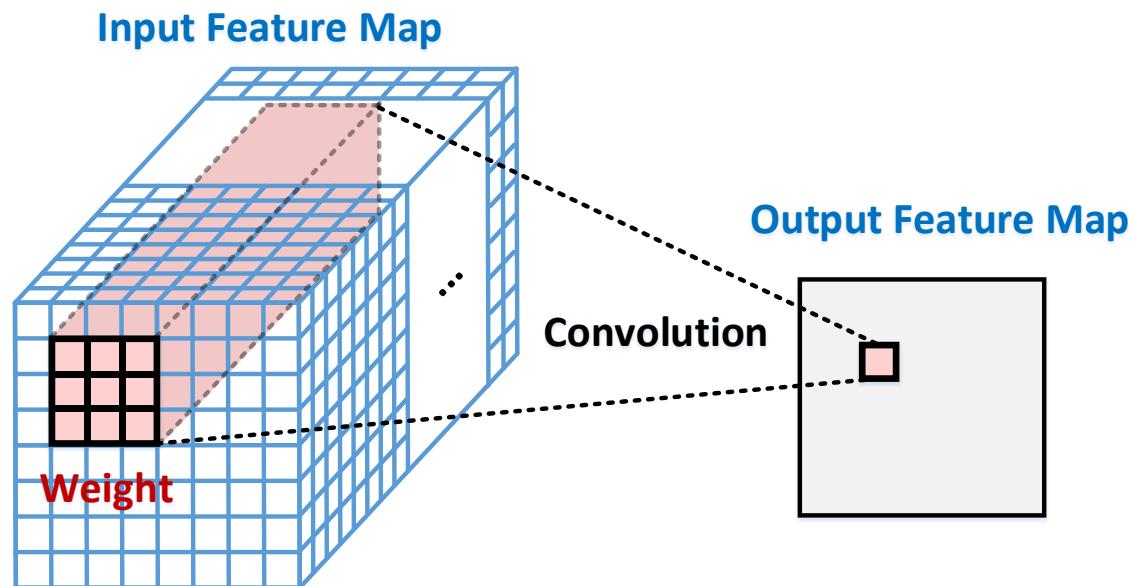
Goal

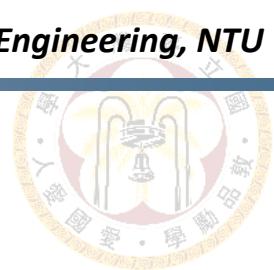
- In this homework, you will learn
 - How to synthesis your design
 - How to run gate-level simulation
 - How to use SRAM



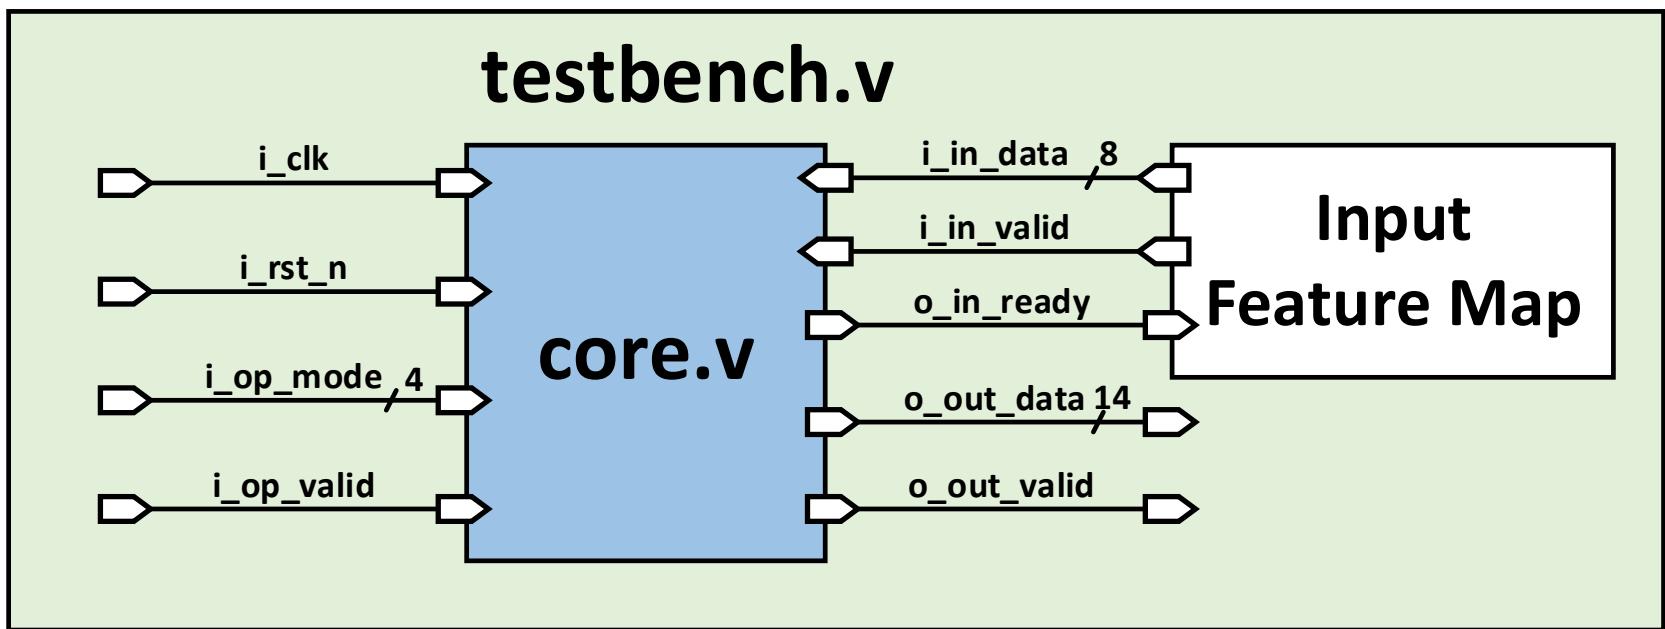
Introduction

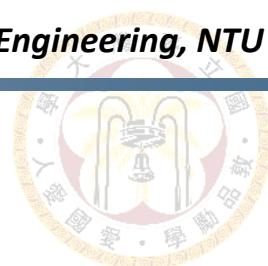
- In this homework, you are going to implement a simplified convolution and image processing engine. An $8 \times 8 \times 32$ feature map will be loaded first, and it will be processed with several functions.





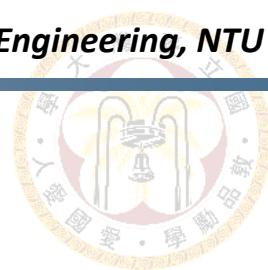
Block Diagram





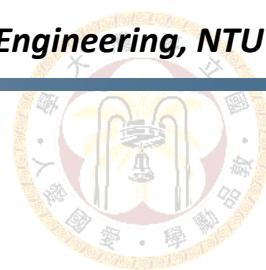
Input/Output

Signal Name	I/O	Width	Simple Description
i_clk	I	1	Clock signal in the system.
i_RST_N	I	1	Active low asynchronous reset.
i_OP_VALID	I	1	This signal is high if operation mode is valid
i_OP_MODE	I	4	Operation mode for processing
o_OP_READY	O	1	Set high if ready to get next operation
i_IN_VALID	I	1	This signal is high if input pixel data is valid
i_IN_DATA	I	8	Input pixel data (unsigned)
o_IN_READY	O	1	Set high if ready to get next input data (only valid for i_OP_MODE = 4'b0000)
o_OUT_VALID	O	1	Set high with valid output data
o_OUT_DATA	O	14	Pixel data or image processing result (signed)



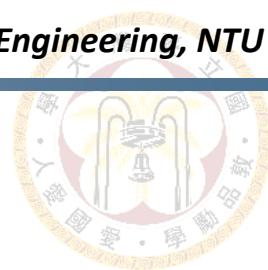
Specification (1)

- All inputs are synchronized with the **negative** edge clock
- All outputs should be synchronized at clock **rising** edge
- You should reset all your outputs when `i_rst_n` is **low**
 - Active low asynchronous reset is used and only once



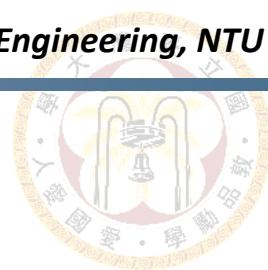
Specification (2)

- Operations are given by i_op_mode when i_op_valid is **high**
- i_op_valid stays only **1** cycle
- i_in_valid and o_out_valid can't be **high** in the same time
- i_op_valid and o_out_valid can't be **high** in the same time
- i_in_valid and o_op_ready can't be **high** in the same time
- i_op_valid and o_op_ready can't be **high** in the same time
- o_op_ready and o_out_valid can't be **high** in the same time



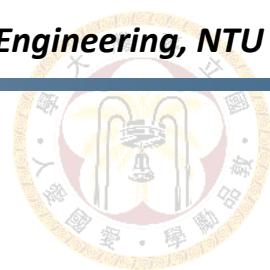
Specification (3)

- Set `o_op_ready` to **high** to get next operation (only one cycle)
 - Raise `o_op_ready` only when the design is prepared for the next operation
- `o_out_valid` should be **high** for valid output results
- **At least one SRAM** is implemented in your design

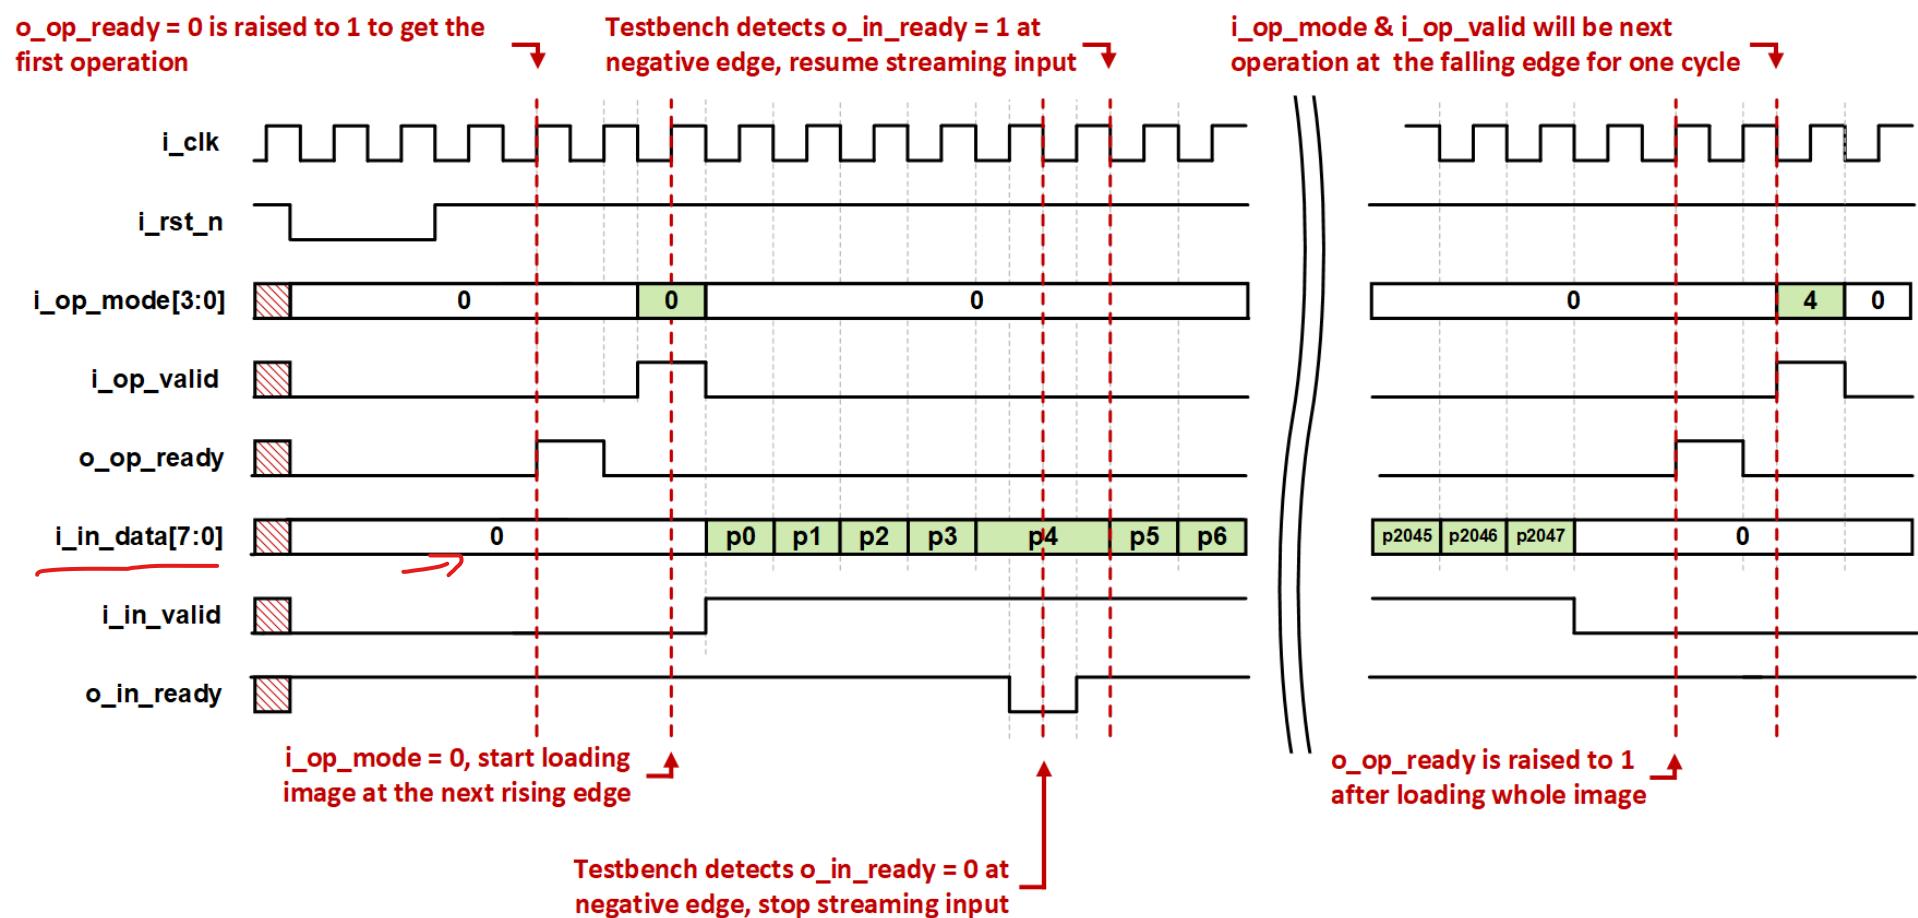


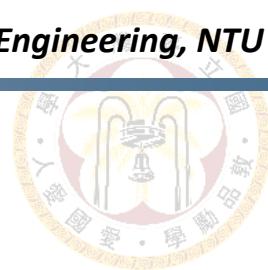
Specification (4)

- Only worst-case library is used for synthesis
- The synthesis result of data type should **NOT** include any **Latch**
- The slack for setup time should be **non-negative**
- **No timing violations or glitches** in the gate-level simulation **after reset**

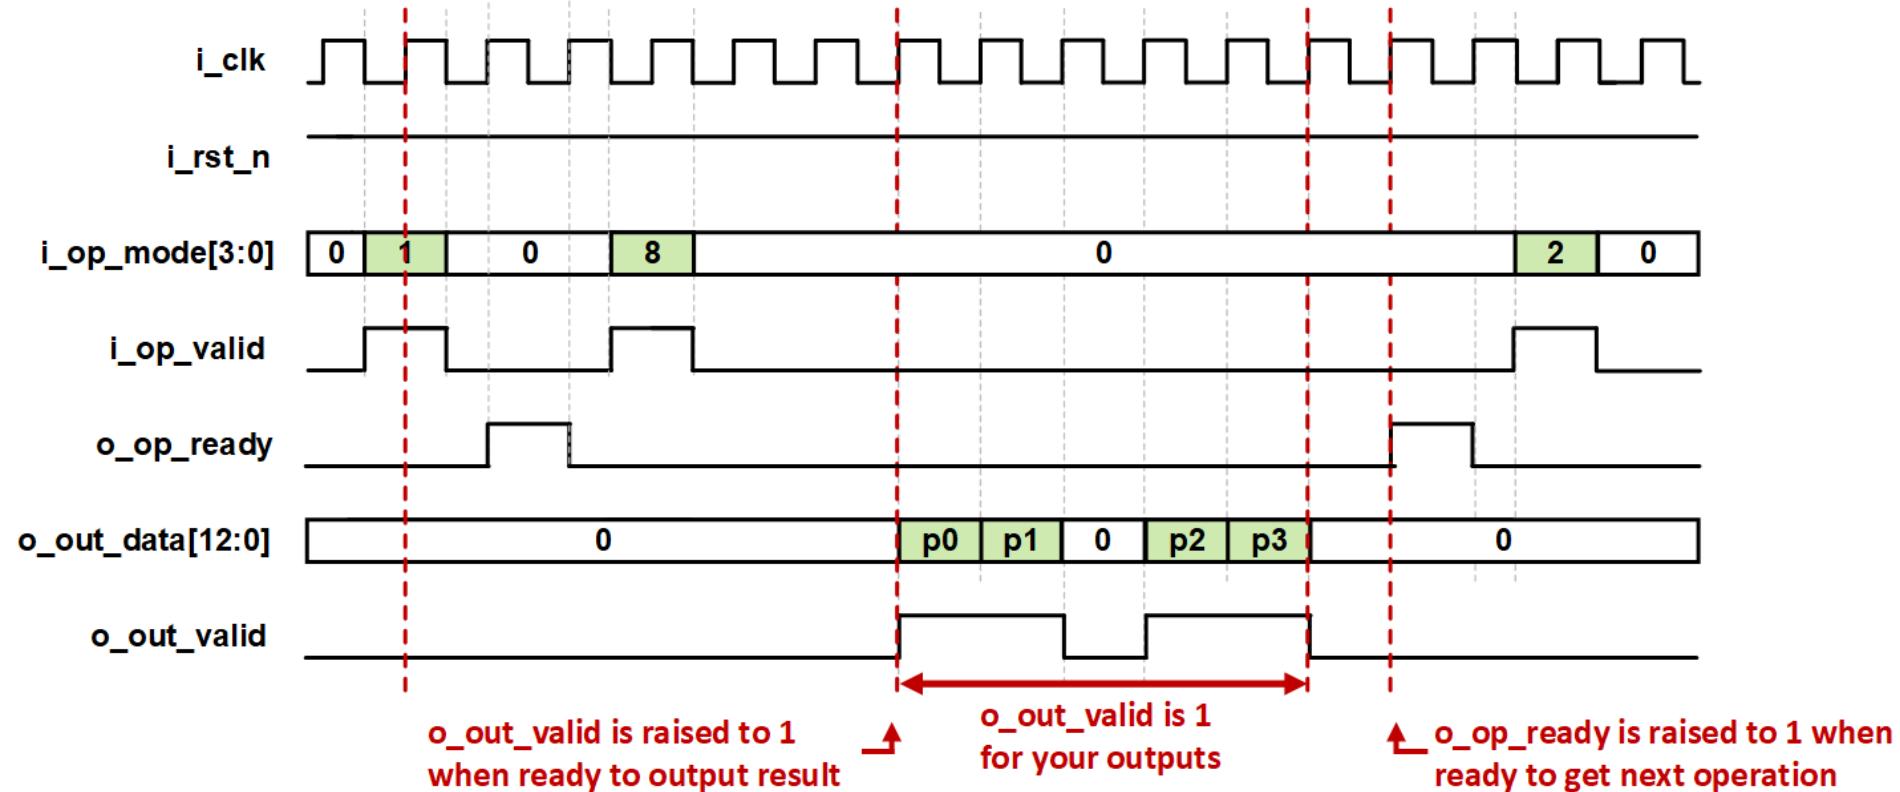


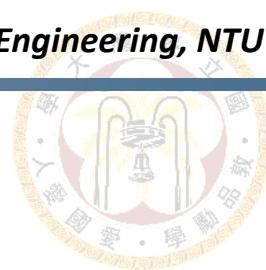
Waveform: Loading Image





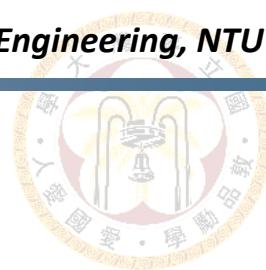
Waveform: Other Operations





Operation Modes

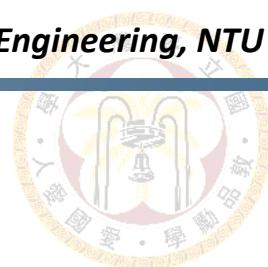
i_op_mode	Meaning
4'b0000	Input feature map loading
4'b0001	Origin right shift
4'b0010	Origin left shift
4'b0011	Origin up shift
4'b0100	Origin down shift
4'b0101	Reduce the channel depth of the display region
4'b0110	Increase the channel depth of the display region
4'b0111	Output the pixels in the display region
4'b1000	Perform convolution in the display region
4'b1001	Median filter operation
4'b1010	Sobel gradient + non-maximum suppression (NMS)



Input Image

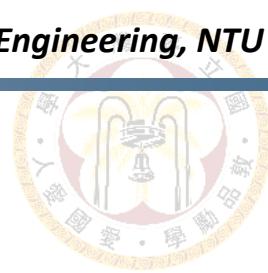
- The input image is given in **raster-scan** order

		channel 31									
		..									
		channel 1									
		channel 0									
		0	1	2	3	4	5	6	7	79	2014 2015
		8	9	10	11	12	13	14	15	87	2022 2023
		16	17	18	19	20	21	22	23	95	2030 2031
		24	25	26	27	28	29	30	31	103	2038 2039
		32	33	34	35	36	37	38	39	111	2046 2047
		40	41	42	43	44	45	46	47	119	..
		48	49	50	51	52	53	54	55	127	..
		56	57	58	59	60	61	62	63		



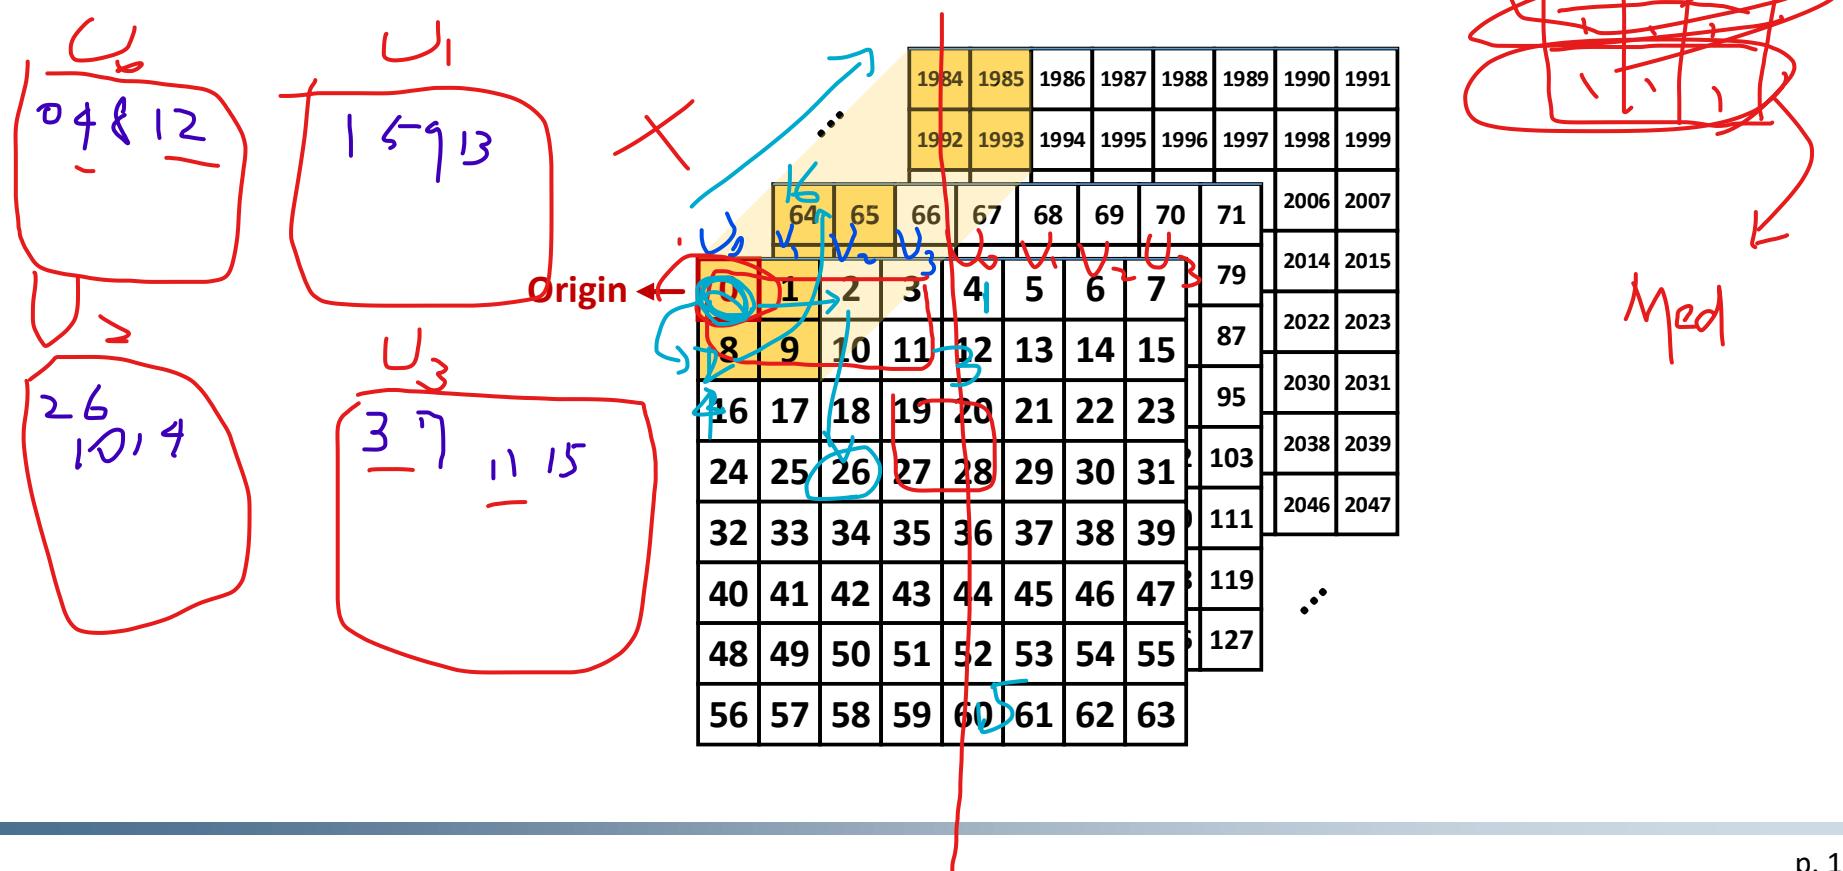
Input Image Loading

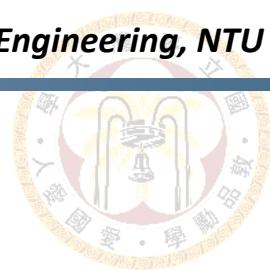
- An 8x8x32 image is loaded for 2048 cycles in **raster-scan** order
- The size of each pixel is 8 bits (unsigned)
- Raise **o_op_ready** to 1 after loading all image pixels
- If **o_in_ready** is 0, stop input data until **o_in_ready** is 1
- The input feature map will be loaded only once at the beginning



Origin

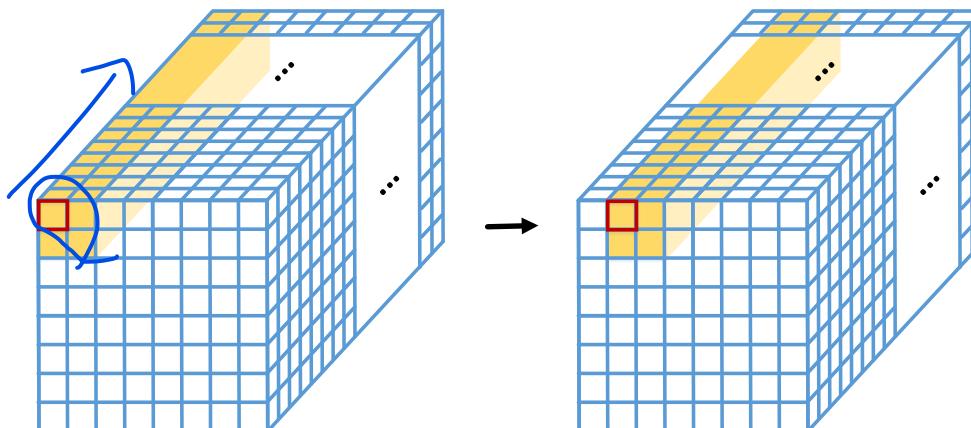
- The first pixel in the display region is **origin**
- The default coordinate of the origin is at 0
- The size of the display region is $2 \times 2 \times \text{depth}$



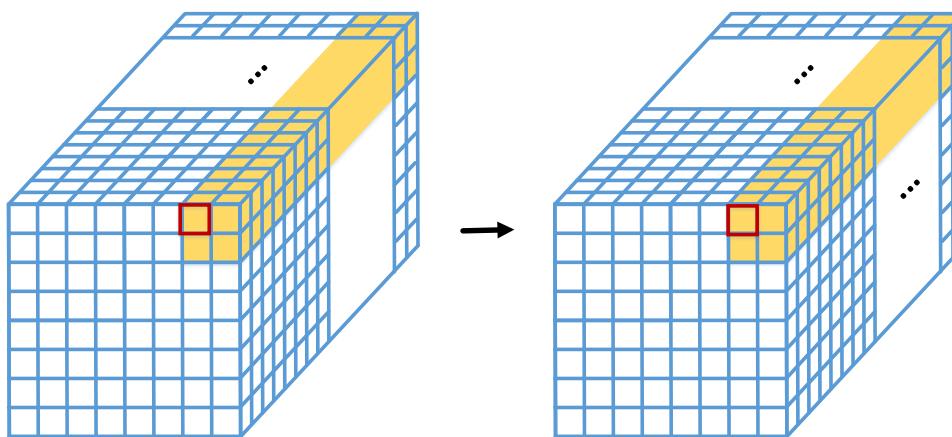


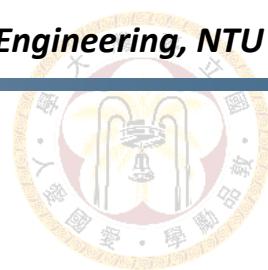
Origin Shifting

- Origin right shift



- If output of display exceeds the boundary, retain the same origin

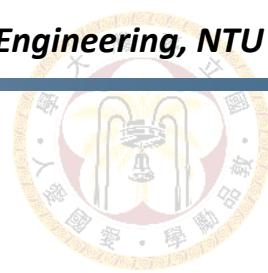




Channel Depth

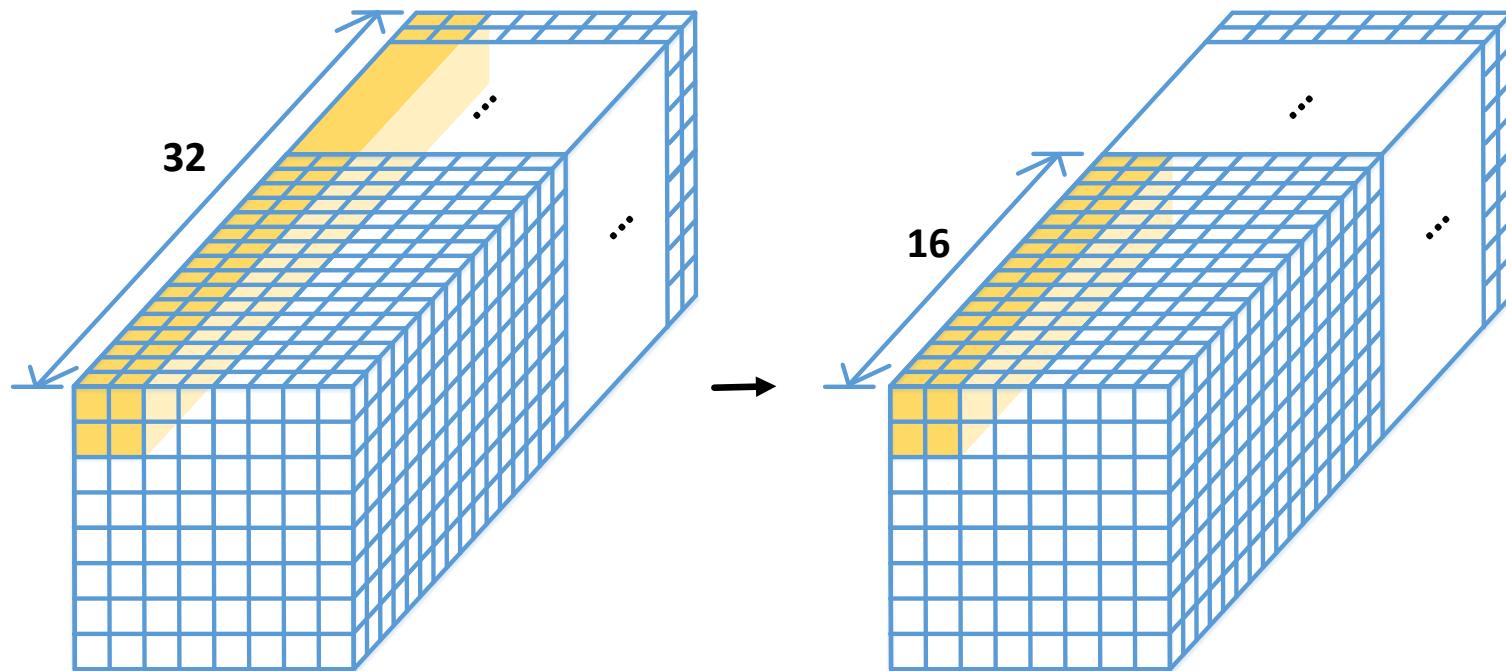
- 3 depths are considered in this design
 - 32, 16, 8
 - Default depth is 32
- The display size will change according to different depth

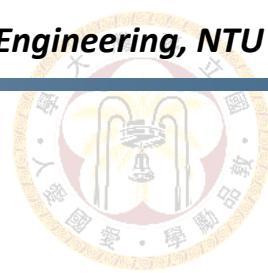
Depth	Display size
32	$2 \times 2 \times 32$
16	$2 \times 2 \times 16$
8	$2 \times 2 \times 8$



Scale-down

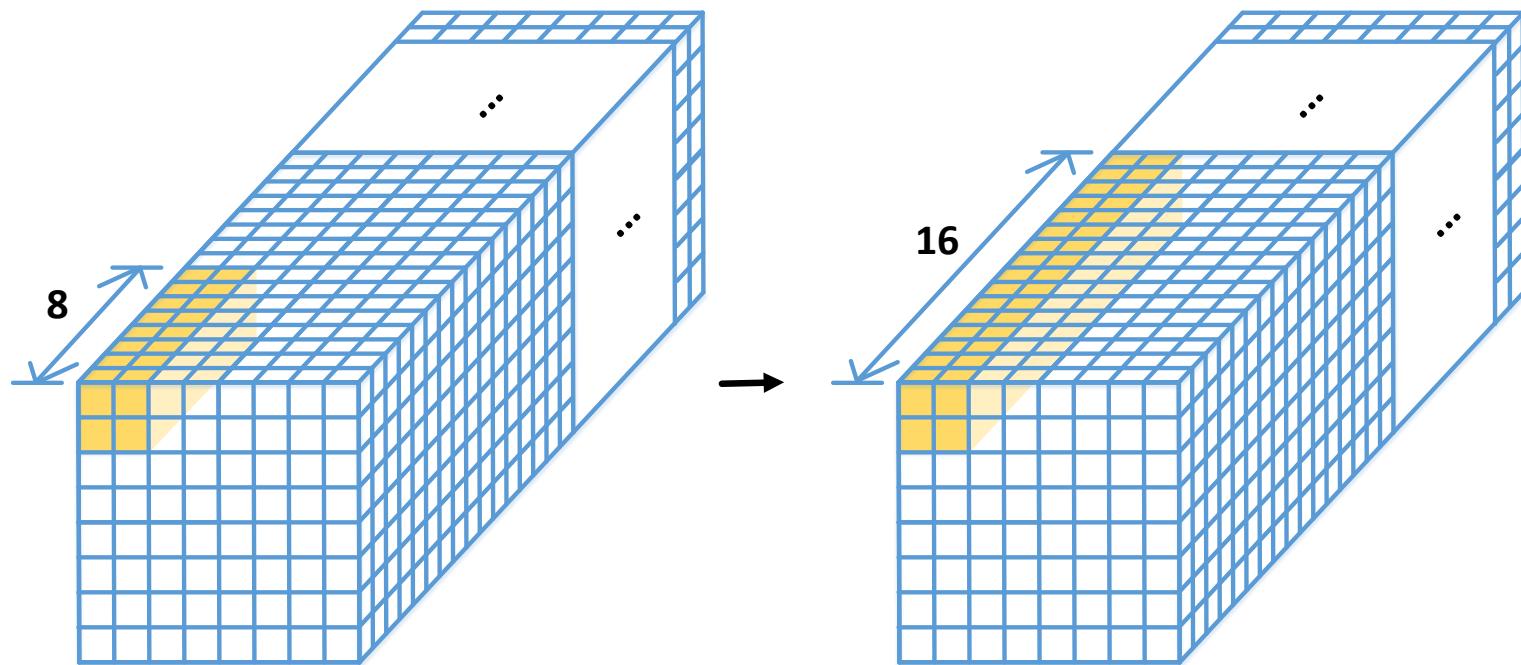
- Reduce the channel depth of the display region
 - Ex. For channel depth, $32 \rightarrow 16 \rightarrow 8$
- If the depth is 8, retain the same depth

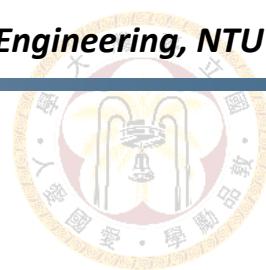




Scale-up

- Increase the channel depth of the display region
 - Ex. For channel depth, $8 \rightarrow 16 \rightarrow 32$
- If the depth is 32, retain the same depth



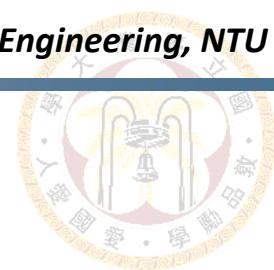


Display

- You have to output the pixels in the display region
- Set **o_out_data [13:8]** to 0 and **o_out_data [7:0]** to pixel data
- The pixels are displayed in **raster-scan** order
 - For example: 0 → 1 → 8 → 9 → 64 → 65 → ... → 1992 → 1993

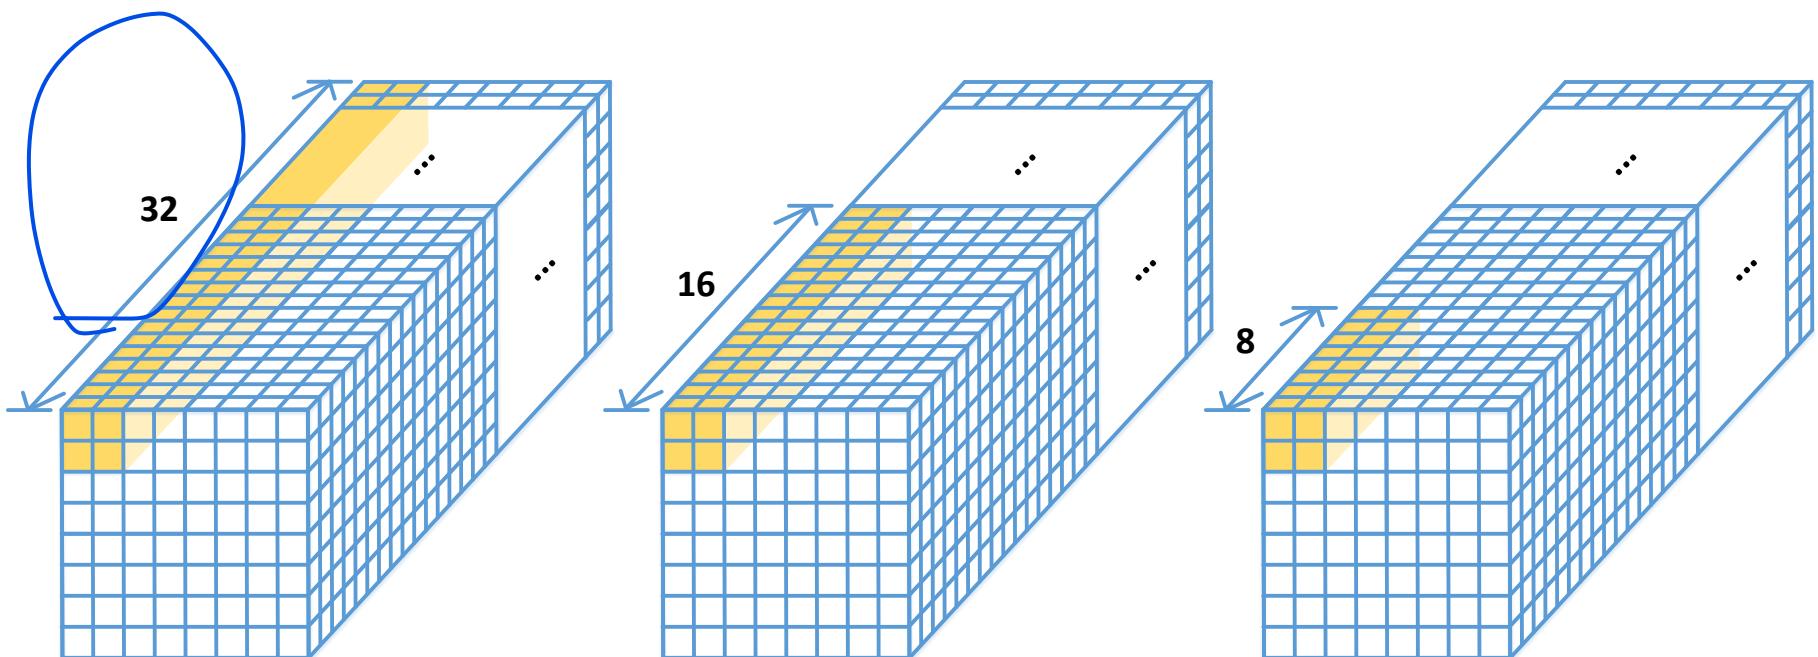
32
16
8

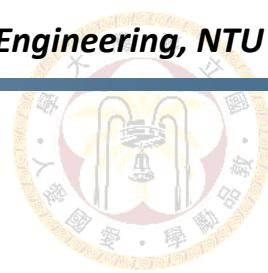
channel 31							
..							
channel 1							
channel 0							
0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63



Display

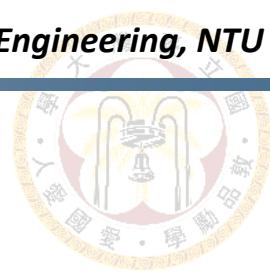
- For display, the display size changes according to the depth



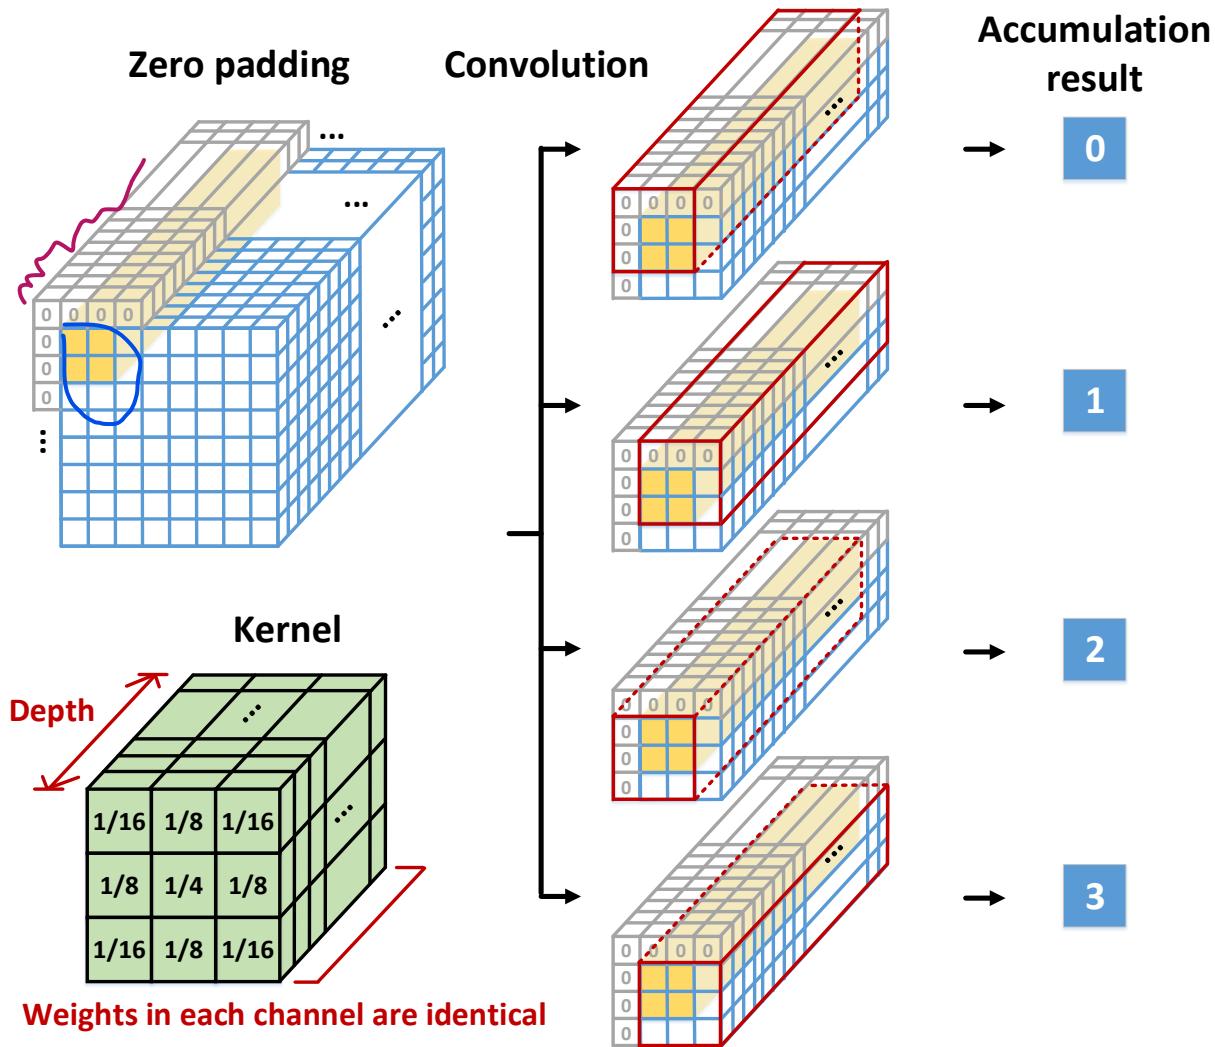


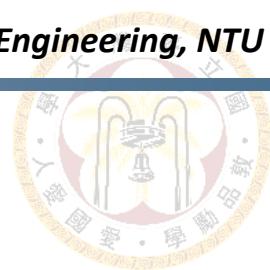
Convolution

- For this operation, you have to perform convolution **in the display region**
- The size of the kernel is a $3 \times 3 \times$ depth and the weights in each channel are identical
- The feature map needs to be zero-padded for convolution
- The accumulation results should be **rounded to the nearest integer** [1]
 - Do not truncate temporary results during computation
- After the convolution, you have to output the **4** accumulation results in **raster-scan** order
- The values of original pixels will not be changed



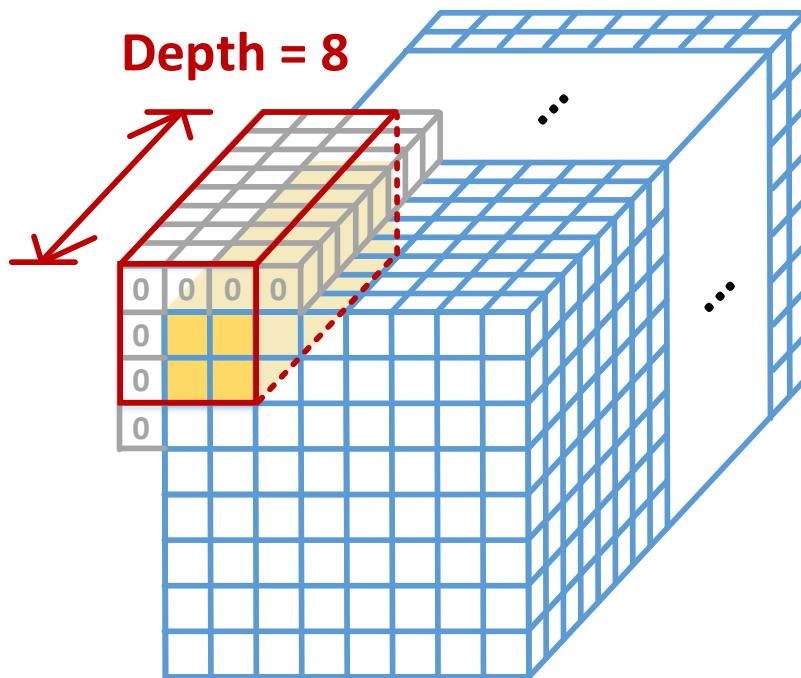
Example of Convolution

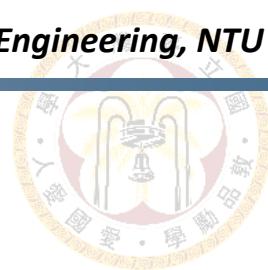




Example of Convolution

- The number of channels that are accumulated during convolution is determined by the depth.
 - For example, accumulate 8 channels if the depth is 8.



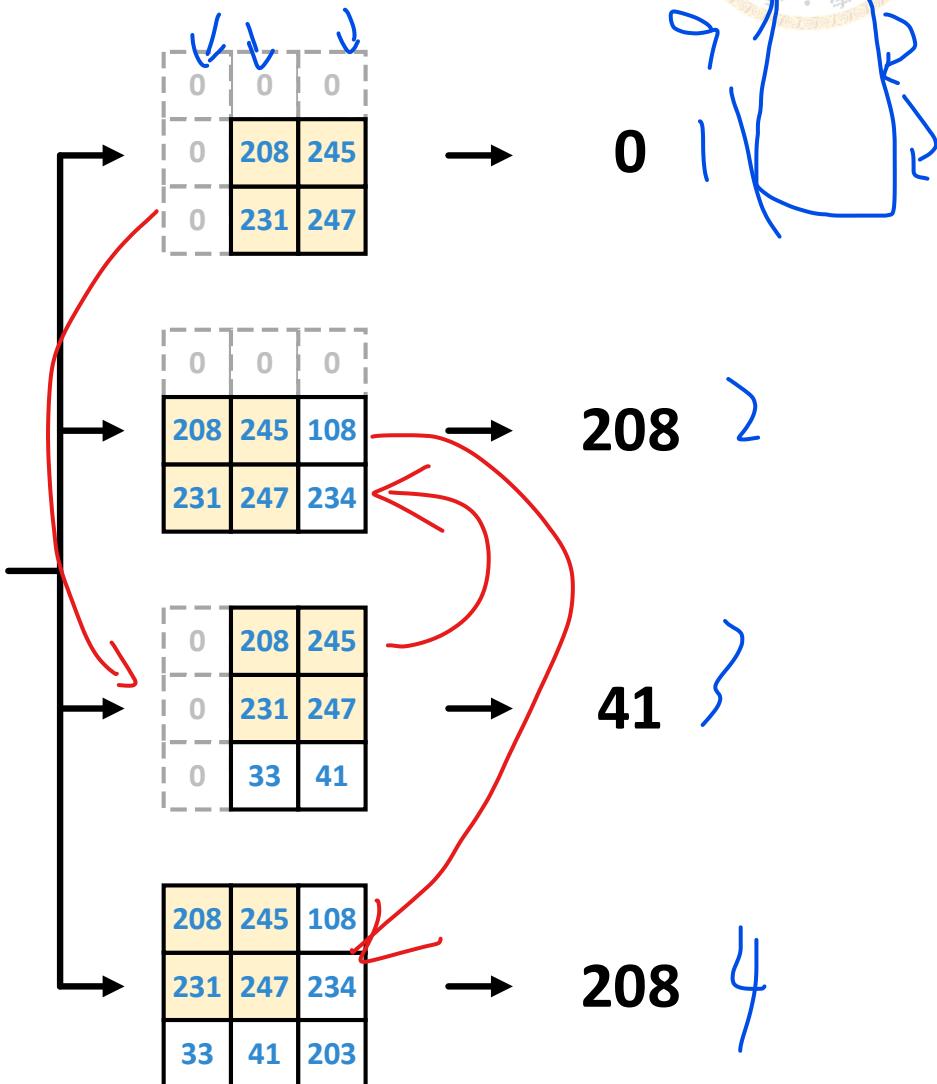
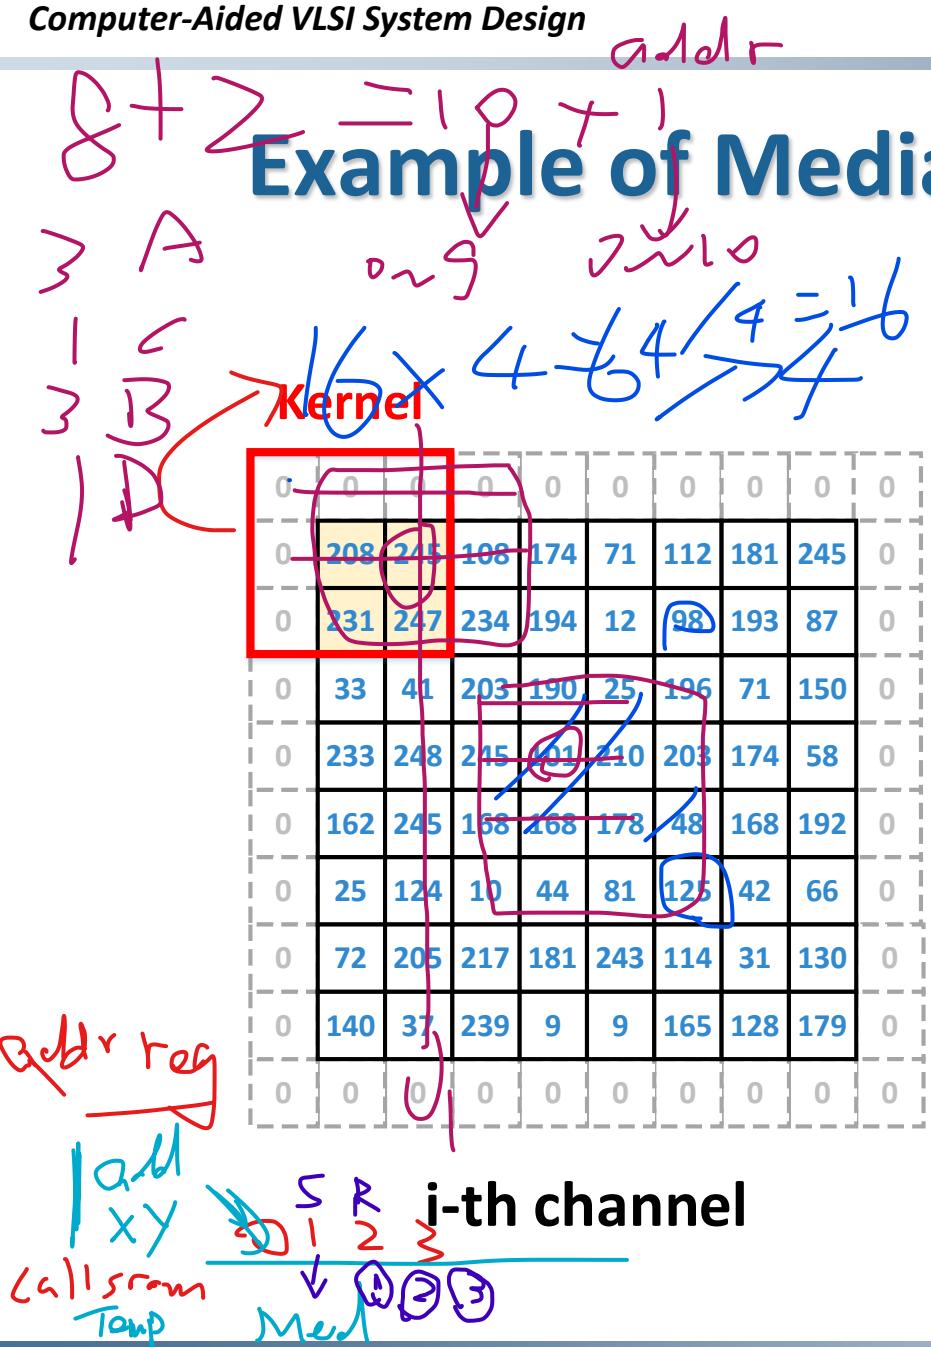


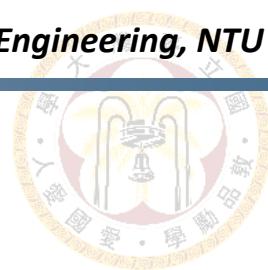
Median Filter Operation

- For this operation, you have to perform median filtering **in the first 4 channels of the display region**
- The kernel size of the median filter is 3×3
- Perform median filtering on each channel **separately**
- The feature map needs to be zero-padded for median filter operation
- After median filtering, you have to output the **$2 \times 2 \times 4$** filtered results in **raster-scan** order
 - Set **$o_out_data[13:8]$** to 0 and **$o_out_data[7:0]$** to pixel data
- The values of original pixels will not be changed

F
chan

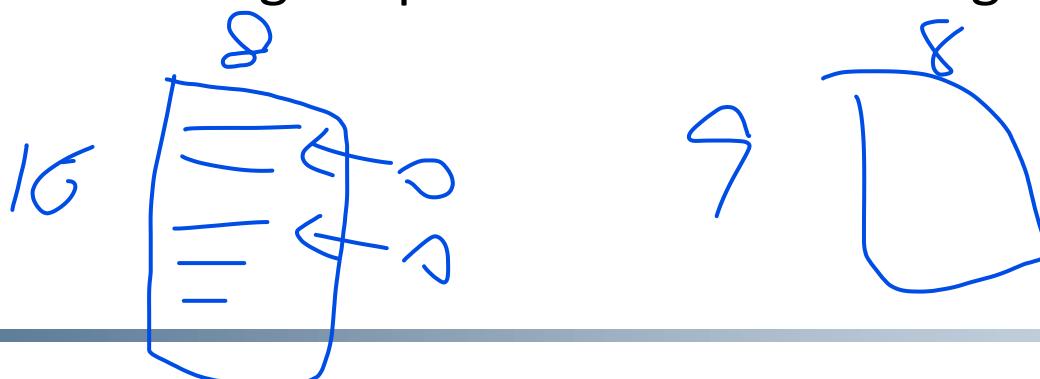
Example of Median Filter Operation

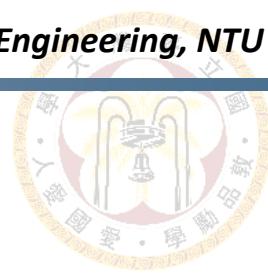




Sobel Gradient + NMS

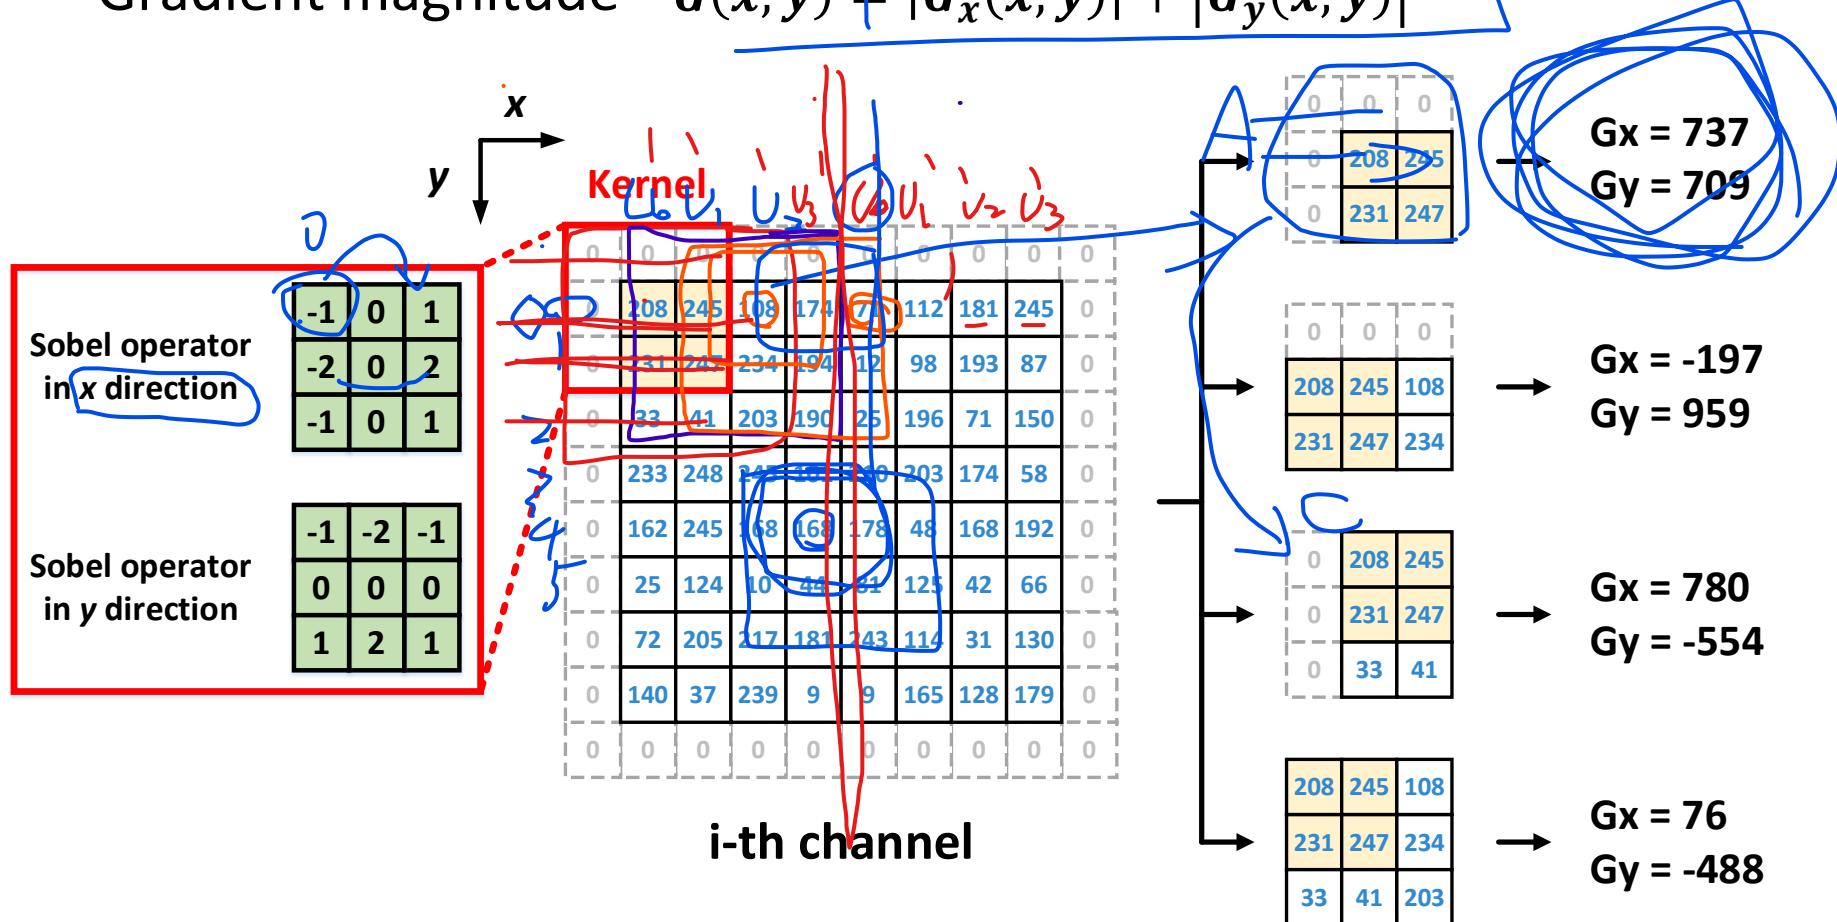
- Calculate gradient **in the first 4 channels of the display region** using the Sobel operator and retain only local maxima along the gradient direction
 - Conduct computations **separately** for each channel
- The kernel size of the Sobel operator is 3×3
- The feature map needs to be zero-padded
- After computation, you have to output the **$2 \times 2 \times 4$** results in **raster-scan** order
- The values of original pixels will not be changed





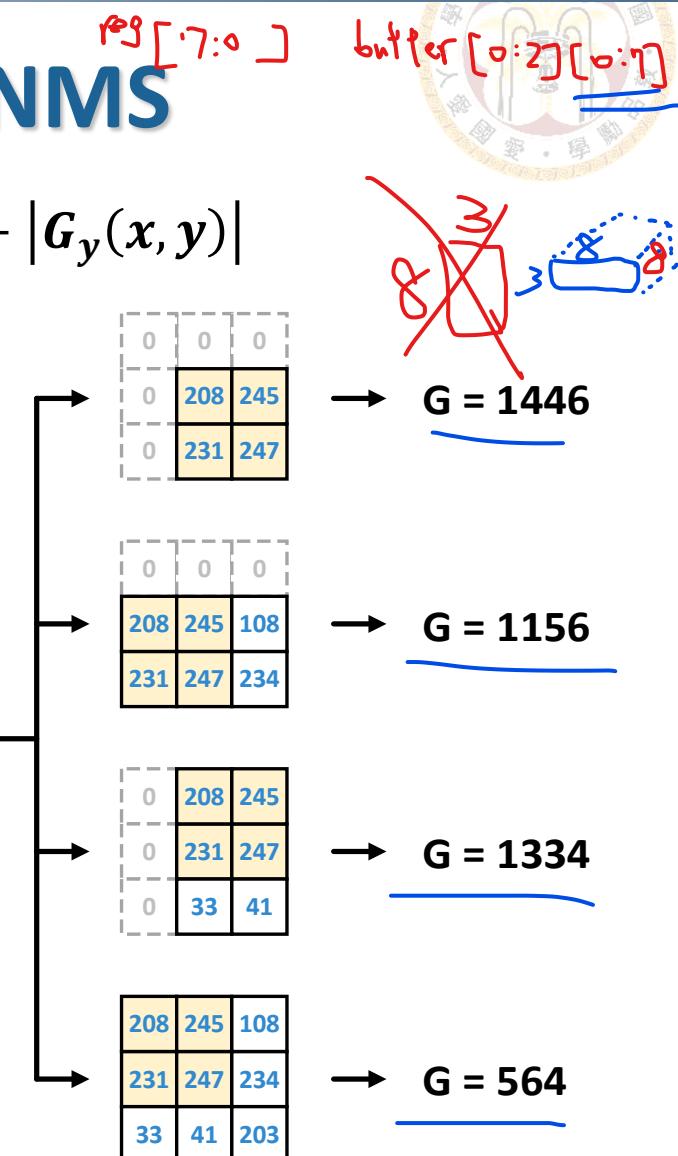
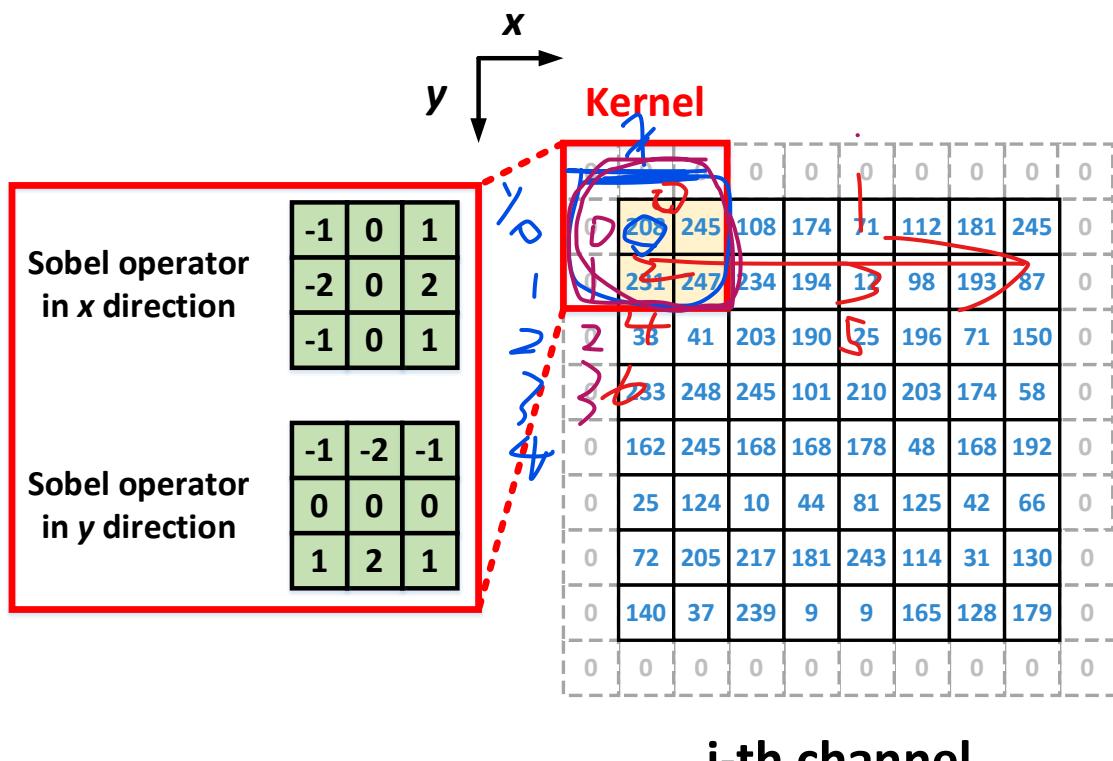
Sobel Gradient + NMS

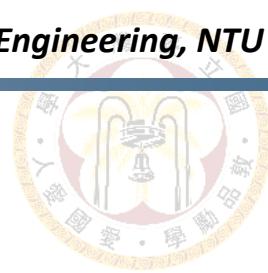
- Gradient magnitude $G(x, y) = |G_x(x, y)| + |G_y(x, y)|$



Sobel Gradient + NMS

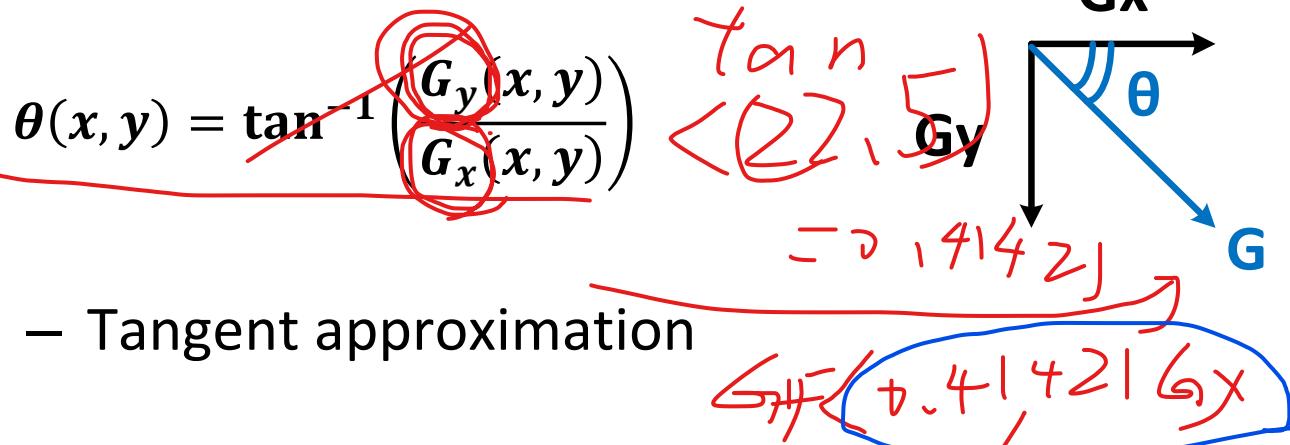
- Gradient magnitude $\underline{G(x,y) = |G_x(x,y)| + |G_y(x,y)|}$





Sobel Gradient + NMS

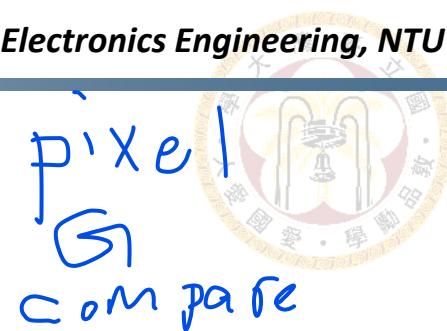
- Gradient direction



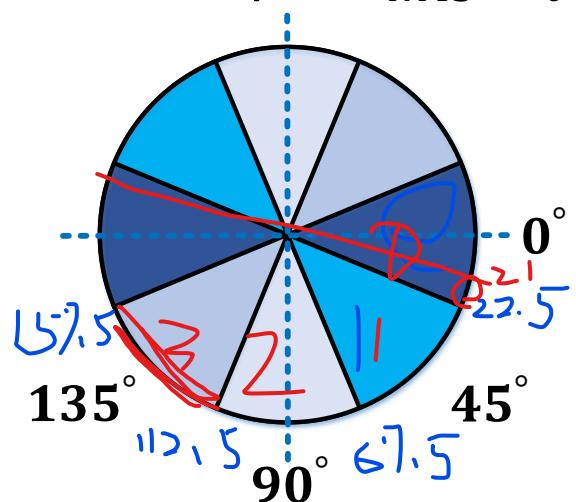
- Tangent approximation

Tangent	Approx. Value	Tangent	Approx. Value
$\tan 0^\circ$	0	$\tan 112.5^\circ$	$-\tan 67.5^\circ$
$\tan 22.5^\circ$	$2^{-2} + 2^{-3} + 2^{-5} + 2^{-7}$	$\tan 135^\circ$	$-\tan 45^\circ$
$\tan 45^\circ$	1	$\tan 157.5^\circ$	$-\tan 22.5^\circ$
$\tan 67.5^\circ$	$2 + 2^{-2} + 2^{-3} + 2^{-5} + 2^{-7}$	$\tan 180^\circ$	$-\tan 0^\circ$

Sobel Gradient + NMS



- Non-maximum suppression (NMS)
 - Find the direction $d_k \in \{0^\circ, 45^\circ, 90^\circ, 135^\circ\}$ that is closest to the gradient direction $\theta(x, y)$
 - If the value of $G(x, y)$ is less than any of its two neighbors along d_k , then set $G_{NMS}(x, y)$ to 0 (suppression); otherwise, set $G_{NMS}(x, y) = G(x, y)$
 - Output $G_{NMS}(x, y)$ in **raster-scan** order

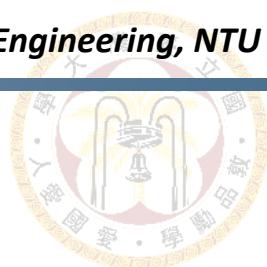


0	0	0	0
0	$G_{(0,0)}$	$G_{(1,0)}$	0
0	$G_{(0,1)}$	$G_{(1,1)}$	0
0	0	0	0

Example:

If $22.5^\circ \leq \theta(0,0) \leq 67.5^\circ$, then compare $G(0,0)$ with its two neighbors along 45° direction, i. e. 0 and $G(1,1)$.

If $112.5^\circ \leq \theta(0,1) \leq 157.5^\circ$, then compare $G(0,1)$ with its two neighbors along 135° direction, i. e. 0 and $G(1,0)$.



Testbench

```
`timescale 1ns/1ps
`define CYCLE      5.0      // CLK period.
`define HCYCLE     (`CYCLE/2)
`define MAX_CYCLE  10000000
`define RST_DELAY   2

`ifdef tb1
  `define INFILE ".../00_TESTBED/PATTERN/indata1.dat"
  `define OPFILE ".../00_TESTBED/PATTERN/opmode1.dat"
  `define GOLDEN ".../00_TESTBED/PATTERN/golden1.dat"
`elsif tb2
  `define INFILE ".../00_TESTBED/PATTERN/indata2.dat"
  `define OPFILE ".../00_TESTBED/PATTERN/opmode2.dat"
  `define GOLDEN ".../00_TESTBED/PATTERN/golden2.dat"
`elsif tb3
  `define INFILE ".../00_TESTBED/PATTERN/indata3.dat"
  `define OPFILE ".../00_TESTBED/PATTERN/opmode3.dat"
  `define GOLDEN ".../00_TESTBED/PATTERN/golden3.dat"
```

```
`elsif tb4
  `define INFILE ".../00_TESTBED/PATTERN/indata4.dat"
  `define OPFILE ".../00_TESTBED/PATTERN/opmode4.dat"
  `define GOLDEN ".../00_TESTBED/PATTERN/golden4.dat"
`else
  `define INFILE ".../00_TESTBED/PATTERN/indata0.dat"
  `define OPFILE ".../00_TESTBED/PATTERN/opmode0.dat"
  `define GOLDEN ".../00_TESTBED/PATTERN/golden0.dat"
`endif

// Modify your sdf file name
`define SDFFILE ".../02_SYN/Netlist/core_syn.sdf"
```

```
// For gate-level simulation only
`ifdef SDF
  initial $sdf_annotate(`SDFFILE, u_core);
  initial #1 $display("SDF File %s were used for this simulation.", `SDFFILE);
`endif
```



Pattern

P20說[13:8]設為0 那golden為什麼[13:8]還有值？

indata*.dat

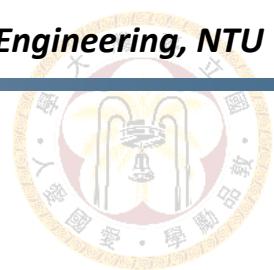
11010000
10101101
10001011
01010110
00101110
11110001
11110110
11001101
01000110
11001110
00001011
10101111
10001000
00111111
00001100
11110100
11100100
00101100
11000100
11001000

opmode*.dat

0000
0010
1000
0100
1000
0011
1000
0100
1000
0110
1000
0001
1000
0110
1000
0011
1000
0100

golden*.dat

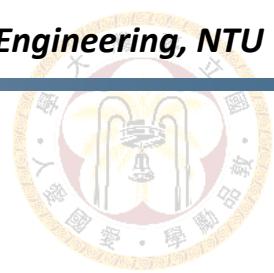
0010011111001
00110011000100
00110001011010
01000001010101
00110001011010
01000001010101
00101110110101
0011111110110
0010011111001
00110011000100
00110001011010
01000001010101
00110001011010
01000001010101
00101110110101
0011111110110
0010011111001
00110011000100
00110001011010
01000001010101
00110001011010
01000001010101
00101110110101
0011111110110



01_RTL

■ core.v

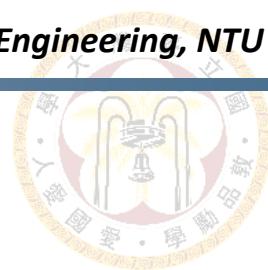
```
module core (                                //Don't modify interface
    input          i_clk,
    input          i_RST_N,
    input          i_OP_VALID,
    input [ 3:0]   i_OP_MODE,
    output         o_OP_READY,
    input          i_IN_VALID,
    input [ 7:0]   i_IN_DATA,
    output         o_IN_READY,
    output         o_OUT_VALID,
    output [13:0]  o_OUT_DATA
);
```



01_RTL

■ rtl_01.f

```
// -----  
// Simulation: HW3  
// -----  
  
// testbench  
// -----  
..../00_TESTBED/testbench.v  
  
// memory file  
// -----  
//.../sram_256x8/sram_256x8.v  
//.../sram_512x8/sram_512x8.v  
//.../sram_4096x8/sram_4096x8.v  
  
// design files  
// -----  
../core.v
```



01_RTL

- Run the RTL simulation under 01_RTL folder

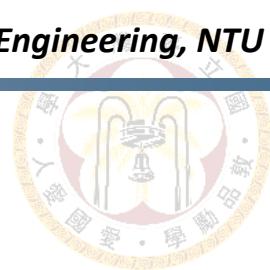
```
vcs -f rtl_01.f -full64 -R -debug_access+all +v2k  
+notimingcheck -sverilog +define+tb0
```

tb0, tb1, tb2, tb3, tb4

or

```
./01_run tb0 5.0
```

clock period



01_RTL

- SpyGlass linting

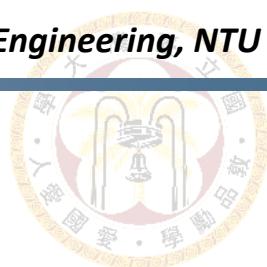
```
./02_lint
```

- Command for cleaning temporary files

```
./99_clean_up
```

- Note that before executing the shell script, change the file permissions by

```
chmod +x ./01_run ./02_lint ./99_clean_up
```



02_SYN

- **core_dc.sdc**

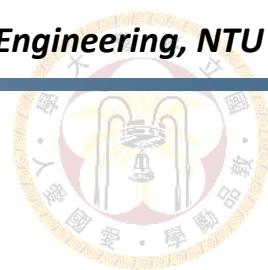
```
# operating conditions and boundary conditions #
set cycle 5.0; # modify your clock cycle here #
```

- **flist.sv**

```
1 // list all paths to your design files
2 `include "../01_RTL/core.v"
```

- Run the command to do synthesis

```
dc_shell-t -f syn.tcl | tee syn.log
```



03_GATE

- Run gate-level simulation under 03_GATE folder

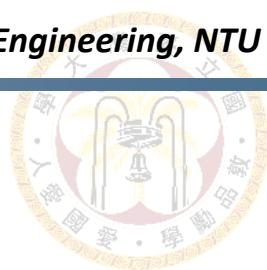
```
vcs -f rtl_03.f -full64 -R -debug_access+all +v2k  
+maxdelays -negdelay +neg_tchk +define+SDF+tb0
```

tb0, tb1, tb2, tb3, tb4

or

```
./03_run tb0 5.0
```

clock period



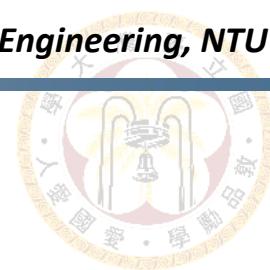
sram_256x8

Pin Description

Pin	Description
A[7:0]	Addresses (A[0] = LSB)
D[7:0]	Data Inputs (D[0] = LSB)
CLK	Clock Input
CEN	Chip Enable
WEN	Write Enable
Q[7:0]	Data Outputs (Q[0] = LSB)

SRAM Logic Table

CEN	WEN	Data Out	Mode	Function
H	X	Last Data	Standby	Address inputs are disabled; data stored in the memory is retained, but the memory cannot be accessed for new reads or writes. Data outputs remain stable.
L	L	Data In	Write	Data on the data input bus D[n-1:0] is written to the memory location specified on the address bus A[m-1:0], and driven through to the data output bus Q[n-1:0].
L	H	SRAM Data	Read	Data on the data output bus Q[n-1:0] is read from the memory location specified on the address bus A[m-1:0].

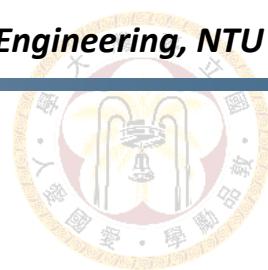


Submission

- Create a folder named **studentID_hw3** and follow the hierarchy below

```
r11943006_hw3/
├── 01_RTL/
│   ├── core.v (and other Verilog files)
│   └── rtl_01.f
├── 02_SYN/
│   ├── syn.tcl
│   ├── flist.sv
│   ├── core_dc.sdc
│   └── Netlist/
│       ├── core_syn.v
│       ├── core_syn.sdf
│       └── core_syn.ddc
└── Report/
    ├── core_syn.area
    └── core_syn.timing
└── 03_GATE/
    └── rtl_03.f
report.txt
```

- Compress the folder **studentID_hw3** in a tar file named **studentID_hw3_vk.tar** (*k* is the number of version, *k* =1,2,...)
 - Use lower case for the letter in your student ID
(Ex. r11943006_hw3_v1.tar)

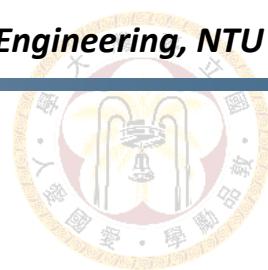


Grading Policy

- Correctness of simulation: **70%** (follow our spec)

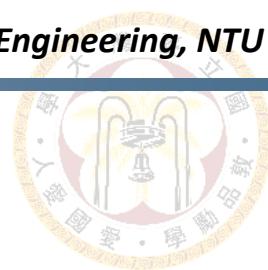
Pattern	Description	RTL simulation	Gate-level simulation
tb0	Load + shift + scale + display	5%	5%
tb1	Load + shift + scale + conv.	5%	10%
tb2	Load + shift + median filter	5%	5%
tb3	Load + shift + Sobel + NMS	5%	10%
tb4	All operations (w/o display)	5%	5%
tbh	<u>Hidden patterns</u>	x	10%

- Performance: **30%**
 - Performance = **Area * Time ($\mu\text{m}^2 * \text{ns}$)**
 - **Time = total simulation time of tb4**
 - The lower the value, the better the performance
 - **Performance score only counts if your design passes all the test patterns**



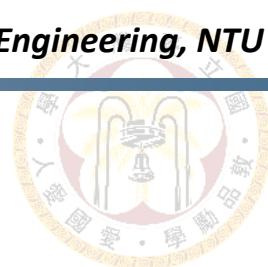
Grading Policy

- **No late submission is allowed**
 - Any submissions after the deadline will receive 0 point
- Lose **5 points** for any wrong naming rule or format for submission
 - Ensure that the submitted files can be decompressed and executed without issues
- **No plagiarism**
 - Plagiarism in any form is strictly prohibited, including copying from online sources or past assignments



Grading Policy

- **Violations of any spec (p.6 – p.9) incur point penalties**
 - Negative slack
 - 0 point for gate-level simulations and performance
 - Design without SRAM
 - 0 point for gate-level simulations and performance
 - Violate other rules but pass all simulations
 - Performance score * 0.7



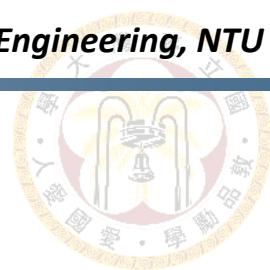
Area

■ core_syn.area

Number of ports:	883
Number of nets:	6074
Number of cells:	5142
Number of combinational cells:	4756
Number of sequential cells:	306
Number of macros/black boxes:	1
Number of buf/inv:	1241
Number of references:	269
Combinational area:	66609.370834
Buf/Inv area:	10982.178019
Noncombinational area:	12156.778545
Macro/Black Box area:	131906.968750
Net Interconnect area:	565696.185242
Total cell area:	210673.118129
Total area:	776369.303371

Number of macros/black boxes
should not be 0

210673.118129 μm^2

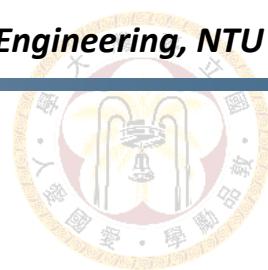


Report

- Your design will be simulated using the clock period in report.txt
- **report.txt**

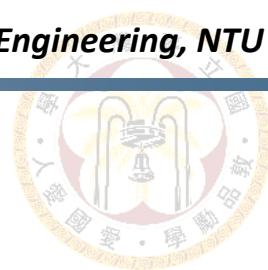
```
1 StudentID:  
2  
3 Clock period: 5.0 (ns)  
4  
5 Area : 210673.118129 (um^2)  
6
```

The clock period that can pass all gate-level simulations without any timing violations



Discussion

- **NTU Cool Discussion Forum**
 - For any questions not related to assignment answers or privacy concerns, please use the NTU Cool discussion forum
 - **TAs will prioritize answering questions on the NTU Cool discussion forum**
- **Email: r11943006@ntu.edu.tw**
 - Title should start with **[CVSD 2024 Fall HW3]**
 - Email with wrong title will be moved to trash automatically



References

- [1] Rounding to the nearest
 - [Rounding - MATLAB & Simulink \(mathworks.com\)](#)
- [2] R. C. Gonzalez and R. E. Woods, *Digital Image Processing*, 4th edition, Pearson, 2018.
- [3] Image gradients and Sobel kernels
 - [Image Gradients with OpenCV \(Sobel and Scharr\)](#)